

Getting started guide of Intel Optane® DC persistent memory module in Kubernetes cluster for Alauda

Author: ailin.yang@intel.com

Homework:

Persistent Memory introduction: [introduction web page](#)

Intel® Optane™ DC Persistent Memory: [Benefit from Greater Capacity, Affordability and Persistence](#)

Concepts of Persistent Memory Provisioning: [Region, Label, Namespace, DAX](#)

Configuring Intel® Optane™ DC Persistent Memory for Best Performance: [configuration video](#)

Provision Intel® Optane™ DC Persistent Memory in Linux*: [video](#)

Configure, Manage, and Profile Intel® Optane™ DC Persistent Memory Modules: [docs](#)

The Intel® Optane™ DC Persistent Memory: [Programming Model](#)

Intel® Optane™ DC Persistent Memory Modules- [Use ipmctl to Debug](#)

Intel® Optane™ DC Persistent Memory Modules – [Provision for KVM/QEMU](#)

System preparation:

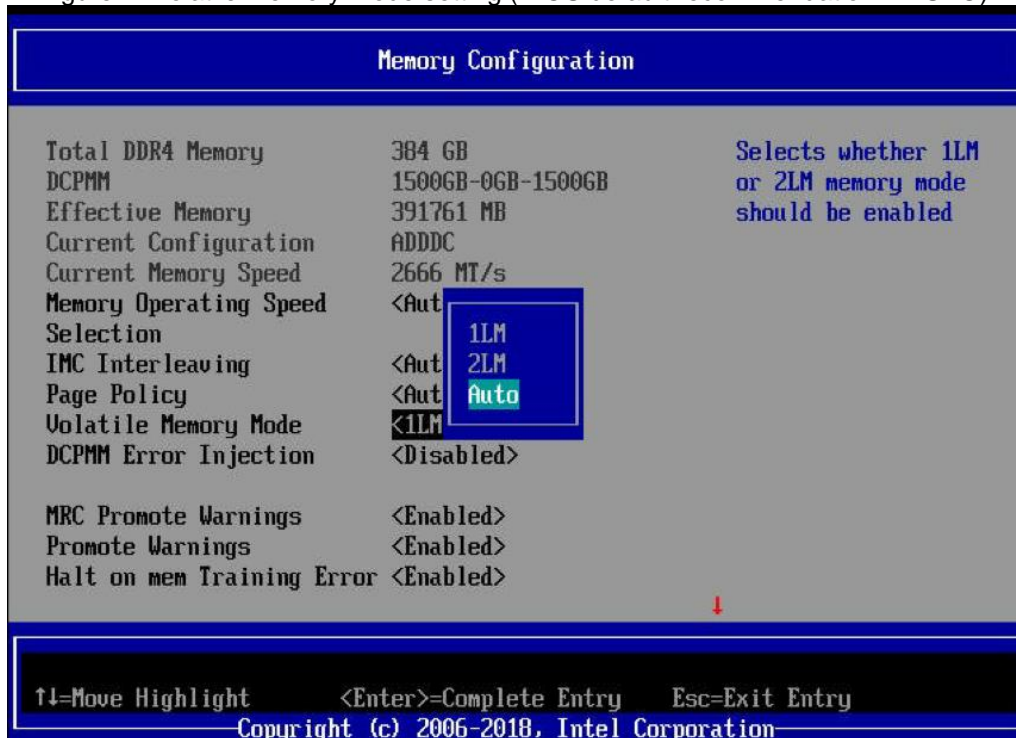
1. 1 Hardware configuration

It depends on the hardware that your server may be different from the one being used in this document. The brief info of the HW is below:

Intel® Server Board S2600WFT, 2x Intel® Xeon® Gold 6252 Processor 2.1GHz 24 cores; 12x 16GB DDR4 DRAM; 12x 128GB Intel Optane DC Persistent Memory

In general, it should be like configure the system to use Intel DCPMM. First it is recommended to confirm that your system has the latest BIOS that support the DCPMM and make sure that Volatile Memory Mode in the Memory Configuration is set to Auto (figure 1).

Figure 1. Volatile memory mode setting (BIOS default recommendation = AUTO).



Note: this is a screen capture of another system with different spec

Figure 2. Topologies per CPU socket (2-2-2, 6 DDR4 DIMMs, 6 DCPMM DIMMs)




Figure 2 shows the memory topology per CPU socket of the reference system. This can be confirmed by the following command:

```
$ sudo ipmctl show -topology
```

DimmID	MemoryType	Capacity	PhysicalID	DeviceLocator
0x0001	Logical Non-Volatile Device	126.4 GiB	0x0028	CPU1_DIMM_A2
0x0011	Logical Non-Volatile Device	126.4 GiB	0x002c	CPU1_DIMM_B2
0x0021	Logical Non-Volatile Device	126.4 GiB	0x0030	CPU1_DIMM_C2
0x0101	Logical Non-Volatile Device	126.4 GiB	0x0036	CPU1_DIMM_D2
0x0111	Logical Non-Volatile Device	126.4 GiB	0x003a	CPU1_DIMM_E2
0x0121	Logical Non-Volatile Device	126.4 GiB	0x003e	CPU1_DIMM_F2
0x1001	Logical Non-Volatile Device	126.4 GiB	0x0044	CPU2_DIMM_A2
0x1011	Logical Non-Volatile Device	126.4 GiB	0x0048	CPU2_DIMM_B2
0x1021	Logical Non-Volatile Device	126.4 GiB	0x004c	CPU2_DIMM_C2
0x1101	Logical Non-Volatile Device	126.4 GiB	0x0052	CPU2_DIMM_D2
0x1111	Logical Non-Volatile Device	126.4 GiB	0x0056	CPU2_DIMM_E2
0x1121	Logical Non-Volatile Device	126.4 GiB	0x005a	CPU2_DIMM_F2
N/A	DDR4	16.0 GiB	0x0026	CPU1_DIMM_A1
N/A	DDR4	16.0 GiB	0x002a	CPU1_DIMM_B1
N/A	DDR4	16.0 GiB	0x002e	CPU1_DIMM_C1
N/A	DDR4	16.0 GiB	0x0034	CPU1_DIMM_D1
N/A	DDR4	16.0 GiB	0x0038	CPU1_DIMM_E1
N/A	DDR4	16.0 GiB	0x003c	CPU1_DIMM_F1
N/A	DDR4	16.0 GiB	0x0042	CPU2_DIMM_A1
N/A	DDR4	16.0 GiB	0x0046	CPU2_DIMM_B1
N/A	DDR4	16.0 GiB	0x004a	CPU2_DIMM_C1
N/A	DDR4	16.0 GiB	0x0050	CPU2_DIMM_D1
N/A	DDR4	16.0 GiB	0x0054	CPU2_DIMM_E1
N/A	DDR4	16.0 GiB	0x0058	CPU2_DIMM_F1

To install ipmctl tool on Ubuntu:

First install the dependencies packages:

```
sudo apt-get install cmake libndctl-dev doxygen build-essential
sudo add-apt-repository ppa:jhli/libsafer
sudo apt-get update
sudo apt-get install libsafer-dev
sudo apt install ruby-full
sudo gem install asciidoctor-pdf --pre
sudo apt-get --no-install-recommends install asciidoc -y
```

Then build and install ipmctl:

```
git clone https://github.com/intel/ipmctl.git
cd ipmctl
mkdir output && cd output
cmake -DRELEASE=ON -DCMAKE_INSTALL_PREFIX=/ ..
make -j all
sudo make install
```

Detail of the tool, please see:

<https://github.com/intel/ipmctl>

<https://github.com/pmem/ndctl>

1.2 Memory mode and App direct mode

1.2.2 Configuring DCPMM in Memory mode

```
$ # destroy all the possible namespaces existing in AppDirect mode
$ # It is necessary to unmount any partitions that are being mounted.
$ # sudo fdisk /dev/pmemXXX -> select "delete" -> select "w"
$ sudo ndctl destroy-namespace all -f
$ sudo ipmctl delete -goal
```

```

$ # create a Memory Mode goal allocation for the entire DCPMM
$ sudo ipmctl create -goal MemoryMode=100
The following configuration will be applied:
SocketID | DimmID | MemorySize | AppDirect1Size | AppDirect2Size
=====
0x0000 | 0x0001 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0011 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0021 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0101 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0111 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0121 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1001 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1011 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1021 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1101 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1111 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1121 | 126.0 GiB | 0.0 GiB | 0.0 GiB
Do you want to continue? [y/n] y
Created following region configuration goal
SocketID | DimmID | MemorySize | AppDirect1Size | AppDirect2Size
=====
0x0000 | 0x0001 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0011 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0021 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0101 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0111 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0000 | 0x0121 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1001 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1011 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1021 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1101 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1111 | 126.0 GiB | 0.0 GiB | 0.0 GiB
0x0001 | 0x1121 | 126.0 GiB | 0.0 GiB | 0.0 GiB
A reboot is required to process new memory allocation goals.
$ sudo reboot
$ sudo ipmctl show -memoryresources
Capacity=1517.1 GiB
MemoryCapacity=1512.0 GiB
AppDirectCapacity=0.0 GiB
UnconfiguredCapacity=0.0 GiB
InaccessibleCapacity=5.1 GiB
ReservedCapacity=0.0 GiB

```

1.2.2 Configuring DCPMM in App Direct mode

```

$ sudo ipmctl delete -goal
$ sudo ipmctl create -goal PersistentMemoryType=AppDirect
The following configuration will be applied:
SocketID | DimmID | MemorySize | AppDirect1Size | AppDirect2Size
=====
0x0000 | 0x0011 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0021 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0001 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0111 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0121 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0101 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1011 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1021 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1001 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1111 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1121 | 0.0 GiB | 126.0 GiB | 0.0 GiB

```

```

0x0001 | 0x1101 | 0.0 GiB | 126.0 GiB | 0.0 GiB
Do you want to continue? [y/n] y
Created following region configuration goal
SocketID | DimmID | MemorySize | AppDirect1Size | AppDirect2Size
=====
0x0000 | 0x0011 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0021 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0001 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0111 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0121 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0000 | 0x0101 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1011 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1021 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1001 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1111 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1121 | 0.0 GiB | 126.0 GiB | 0.0 GiB
0x0001 | 0x1101 | 0.0 GiB | 126.0 GiB | 0.0 GiB
A reboot is required to process new memory allocation goals.
$ sudo reboot

```

Deploy Pmem-csi in Kubernetes cluster

In this documentation, the deployment is based on an exist Kubernetes cluster, so setup a Kubernetes cluster is not covered, and Before start, there is something that needs to be careful here. First please make sure that the system has been configured to use Intel DCPMM in App Direct mode. Second, it is necessary to make sure that the volume groups are clean. It can be checked by using the following command.

```
$ sudo vgs
```

Please confirm that it should return nothing before we move forward.

Also please read readme of <https://github.com/intel/pmem-csi> to know requirements to deploy it,

Now you can follow up [this screencast guide](#) to deploy Pmem-CSI in your Kubernetes cluster

After completing deployment, you should have it run correctly with below status:

```

PMEM-CSI $ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
my-csi-app-1                        1/1     Running   0           8s
my-csi-app-2                        1/1     Running   0           8s
pmem-csi-controller-0              2/2     Running   0          2m10s
pmem-csi-node-jrfck                2/2     Running   0          2m10s

```

Congratulations! You have PEME-CSI ready